

AI-Powered Real-Time Log Monitoring and Incident Response System for Web Applications

Mrs. S. Ruby¹, T. V. Akhil², G. Akshay³, J. Godwin Judah⁴, Mr. T. Pandiarajan⁵

^{1 5}Assistant Professor, ^{2 3 4}Undergraduate Student, Department of Computer Science Engineering,
Rajalakshmi Institute of Technology, Chennai, India

ABSTRACT

Due to distributed systems, microservices, and continuous user traffic, modern web applications generate enormous amounts of log data, which leads to slow, imprecise, and ineffective log analysis techniques. Long downtime and low system reliability can result from late detection of critical events. The given paper suggests a real-time log monitoring and incident response system powered by AI implemented to monitor the web applications that will automate the process of log ingestion, analysis, and anomaly detection. The system applies anomaly detection methods that are anchored on machine learning and matches events among various application components. Through analyzing live log streams, the framework can identify possible failures in time, and alerts and remediation information can be taken through an interactive dashboard. The suggested solution lowers operational overhead, quickens incident reaction, and enhances general versatility of the applications with scalable and open-source innovations.

KEYWORDS

Monitoring of logs, detection of anomalies, incident response, web application, machine learning, real time analysis, event correlation, AIOps.

I. INTRODUCTION

Modern web apps play a significant role in the delivery of reliable digital services, but they constantly generate enormous volumes of log data on servers, databases, microservices, and APIs. These logs have useful data regarding system behavior, performance and failure. However, manual log verification is inefficient, time-consuming, and very error-prone, especially in large-scale distributed environments. As applications become more complex, prompt detection of errors and anomalies may result in downtime delays, poor user experience, and increased operating expenses

Conventional log monitoring and management platforms are deeply entrenched in fixed rules, alerts which are based on thresholds or hand-coded configuration. These methods are effective at simulating known issues, but they are unable to adapt dynamically to shifting workloads, system behavior, or novel failure behaviors. In addition, they have no means to correlate intelligently the events between multiple components and therefore, root-cause analysis is slow and challenging. Because of this, system administrators tend to respond to incidents when there is a lot of damage done.

Due to the development of machine learning and AIOps, intelligent and automated log examination has become a potentially attractive approach to proactive monitoring of the system. The machine learning technologies are able to predict the normal behavior of a system based on the past logs and

anomalies in real-time without having to depend on numerous manual rules to define. The techniques can help with early critical incident detection, improved incident correlation, and prompt incident response.

The current paper proposes a web application-based REAL-time log monitoring and incident response system that will operate using AI. The provided system uses machine learning to automate log ingestion, preprocessing, and anomaly detection. It processes application log streams in real-time to detect the presence of an unusual pattern, cross-correlates events among application components, and provides actionable insights as an interactive dashboard. The system is set to be scalable, efficient and open-source based technologies can be deployed, which is why they enhance application reliability and minimize operational overhead.

II. LITERATURE REVIEW

Numerous disciplines, including system observability, distributed systems, machine learning, anomaly detection, and AIOps, have studied log monitoring and incident detection. This part will review current methods of analysis of logs, discuss their weaknesses, and outline the research gaps being filled with the proposed AI-based real-time log monitoring and incident response system.

A. Conventional Log Management and Monitoring Systems.

ELK Stack, Splunk, and Graylog are examples of traditional log monitoring systems that provide centralized log aggregation, indexing, and visualization. These systems allow an administrator to search and analyze predefined queries and dashboards. Despite this, they still rely

largely revolving around manually set rules, thresholds, and notifications that need, “domain knowledge.” Those rule-based methods find it hard to fit dynamic behavior of a system and tend not to identify the failure patterns that have never been observed before or hidden anomalies in large-scale distributed settings.

B. Incident Detection and Rule-Based Systems Automated.

A number of automated monitoring tools use constant thresholds or signature based guidelines to identify incidences. Although efficient in detecting well-known faults, crashes or resource saturation cases, the techniques are not capable of responding to complex and dynamic workloads. They tend to produce high false-positive rates when workload spikes and can not correlate the related events across services. With the shift to microservices and cloud-native infrastructure in modern applications, rule-based systems are becoming less helpful than they were once in real-time incident response due to their inflexibility.

C. Log Anomaly Detection by machine learning.

The use of machine learning methods has become eminent in the field of log analysis since it has been viewed to be able to model the normal operation of the system and also detect any deviations as an automatic system. Methods that use PCA to calculate the variance between records have been popular in the context of detecting anomalies in logs, as well as methods of clustering and unsupervised learning tools like Isolation Forests. More recent works use deep learning models to extract temporal and sequential patterns of log streams. Although they enhance detection accuracy, most of them are designed to provide offline analysis or isolated components and do not effectively aid real-time detection and multi-component event correlation.

D. Event Correlation and Root-Cause Analysis.

The event correlation methods are used to establish the association between the log events produced by various services and system layers. Correlation of logs is essential in learning about cascading failures and root causes of distributed systems. The correlation techniques that are currently in use usually involve heuristics, timestamps or dependencies among services, which might not be adequate in high concurrent environments. Inadequate contextual awareness and inbuilt anomaly detection make them less effective to offer actionable insights in live incidents.

E. Research Gap and Contribution of the Current Work.

The existing solutions for log monitoring are beset with serious drawbacks such as being too dependent on human rules, a lack adaptation to changing systems, inability to provide real-time support, and poor inter-service correlation. Most machine learning systems are not well integrated with incident response processes and lack visualization. In order to close such knowledge gaps, the system that is going to be proposed comprises a number of real-time log ingestion, machine learning anomaly detection, and intelligent event correlation into a single system. The system will offer early warning of incidents, contextual understanding, and remediation assistance through AI-based analysis and an interactive dashboard, allowing to resolve the incident faster and enhance the reliability of the application.

III. PROPOSED METHODOLOGY

The suggested system incorporates the end to end intelligent log monitoring system and incident response model that will support the new web applications. The system is also data-driven, automated and can analyze in real time as compared to the old system of analyzing logs, which involves application of fixed rules or human inspection of the logs. It employs machine learning to detect anomalies in the log stream and correlate events in the parts of the distributed application. The solution brings the log ingestion, preprocessing, feature extraction, anomaly detection and anomaly correlation together into a single pipeline. The logs that are collected through the API and backend server are continuously stored and made in the machine-readable formats. Relevant attributes such as timestamps, log frequency, error codes and service identities are obtained to model normal operations of the system. This is then run through an anomaly detecting engine which examines incoming logs in real time in order to identify deviations that may indicate failures or deterioration in performance. These broad steps encompass five major processes, namely the log data acquisition process, the process of preprocessing and normalization of the log data, features extraction, machine learning-based anomaly detection, and event-correlated alert creation. The system is continuously watching live log streams and alerts and introduces insights to a monitoring dashboard in a timely manner. It is an automated, real-time approach, which enables the early detection of incidences, detect and diagnose it faster and offer more reliability to the efforts in the complex web application environment.

IV. SYSTEM ARCHITECTURE

Artificial Intelligence-based Log monitoring system architecture.

System architecture is based on a layered and modular design to assist in the effective log ingestion, real-time processing, anomaly detection, and incident response. It is intended to process high volume log streams produced by the modern web applications and deliver timely and actionable insights. The architecture is based on a centralized monitoring interface, Python-based log processing and machine learning backend,




Figure 1. Methodology Overview

and a data storage layer that can be used to analyze both in real-time and historical.




Figure 2. System Architecture

A. Architectural Components

- The architecture has various functional layers.
- The Log Collection Layer takes logs off web servers, application services, and APIs and the backend components.
- Log Processing Layer Parses, cleans and organizes the raw log messages into a normalized format to be analyzed.
- The Anomaly Detection Layer uses machine learning models to detect abnormal behavior of the system.
- The Event Correlation and Alerting Layer works with anomalies between services and creates alerts that include context-related information.
- The Visualization and Dashboard Layer displays real-time information, anomaly summaries, and incident reports to the system administrators.
- This is a modular architecture that provides scalability, flexibility, and maintenance as well as real-time monitoring.

B. Preprocessing Layer

The preprocessing layer is the layer that performs pre-processing of the raw log data so that the incoming log messages are cleaned and normalized so that they can be analyzed. It does normalization of timestamps, deletion of unwanted fields, log parsing and standardization of features to support differences in log format across services. Methods applied to reduce redundant information, and noise filtering, to make sure that important system events are not lost. This is a preprocessing step that enhances accuracy and reliability of downstream detection models of anomalies.

C. Data Storage Layer

The system uses a combined strategy of data storage to control real time and historical log information. Short term storage keeps the recent logs and extracted features that are useful in the detection of anomalies and creation of alerts very quickly. Storage keeps historical records, anomaly records, and incident metadata to carry out the trend analysis, auditing, and model retraining. This will make the system scalable and allow more thorough investigation of system behavior as time goes by.

D. Machine Learning and Anomaly Detection Layer.

The machine learning layer is the heart and soul of the system and it is based on historical data to model normal log behavior. Frequency, pattern of errors, service identifiers, and time aspects are analyzed to segregate abnormalities in real-time. Unsupervised learning methods are utilized when unusual patterns are determined with no labeled information on their failure. The detected ones are rated according to the degree of anomaly and then sent to these triggers are then fed into a correlation engine for further processing.

E. Monitoring Interface and Alerting Layer.

The monitoring interface offers an easy to use dashboard which shows detected anomalies, correlated events as well as system health metrics. Alerts occur when there are critical anomalies and are further classified based on severity, to help in the quick decision making. The interface allows operators to view detailed log traces, learn about the impacted services and take remedial actions as they come.

F. Architectural Advantages

The proposed architecture introduces several benefits in terms of real- time monitoring log. It does not depend on a fixed set of rules. In this scenario, it becomes highly relevant because it applies machine learning models which are adaptive in anomaly detection is the process of identifying data points in a dataset that differ significantly from the remaining data. Processing in real-time allows the identification early warning of incidents and rapid responses. The scalability of modular architecture allows the scaling up of a large distributed systems and the scalability of the architecture to integrate with other monitoring tools to be easily integrated into the system. The system also increases effectiveness when root-cause analysis by correlating various events in several components, thereby making web applications more robust and resilient.

V. EXPERIMENTAL SETUP

This section outlines the dataset, experimental setup, system set up, and evaluation process that was to be employed to determine the performance of the proposed AI-based real time log monitor and incident response system. Because there is no single public dataset which comprehensively represents the actual-world web application

logs of multiple components and failure conditions, a custom dataset was formed by merging publicly available log datasets with synthetically created and simulated log events.

A. Dataset Construction

In order to train and test the system, a multi-source log dataset was developed to match real-world application of web applications. The baseline to record normal operation behavior was the publicly available system and applications log datasets. Moreover, the artificial record of the logs was created to emulate the typical failure conditions, including the crashing of the service, request time outage, database connection, and the development of the abnormal rates of the




Figure 3. System Workflow

traffic. To capture a distributed environment, logs were taken on various components such as web servers, application services and backend modules. Artificial anomalies were also introduced to test both the performance of the anomaly detector and event correlation algorithms. These were abrupt error rate, abnormal response delay, work restart, and lost workflow events. The resultant dataset will have both regular and anomalous log patterns which can be used to train unsupervised machine learning models and test actual real-time detection services.

B. Composition and Characteristics of the Dataset.

The developed dataset provides a variety of system behaviors that one may meet in real-life implementations. It contains informational logs of high frequency, warning and error messages, timestamped service calls and structured and semi-structured log formats. Normal logs are used to represent the actual operation states by ensuring that the data is homogeneous, and the anomalous logs are distributed sparsely to make sure that the system is sensitive to some rare events. The data is intended to have noise, such as variations in log format, incomplete messages, as well as bursty traffic patterns to production environments. The variability inherent in such a system makes capable of generalizing from workloads, application architectures, to

how cutting is to be performed. It is possible to use realistic operation data, and injected anomalies, which serve as a strong base on which to examine the precision of detection, scalability, real-time performance.

VI. RESULTS AND DISCUSSION

As the experimental analysis of the suggested AI-driven log monitoring and incident response system proves, it works well in terms of real-time anomaly detection, event correlation, and incident identification. The anomaly detection module was

able to detect an abnormal pattern in logs that are related to service failures, lowered performance and unforeseen execution behavior despite the noisy and high volume log streams. The fact that the system has a high detection accuracy and low false-positive rates shows that the system is reliable when differentiating critical incidents and normal operational variations.

The event correlation mechanism successfully clustered similar anomalies across various application elements which were found to be quicker in identifying the root causes in the distributed environment. The proposed system demonstrated a higher degree of adaptability to dynamic workloads and requirement in fewer manual-defined rules in comparison with the traditional rule-based log monitoring approaches. The system could identify failure conditions never seen before which is common in the case of the old style of threshold-based approaches.

The built-in monitoring dashboard was able to give clear visualization of the anomalies, timeline of incidents and services affected, which increased situational awareness to operators. The qualitative observation of observed incidents ensured that the system was suitable to record cascading failures, repeat error patterns as well as abnormal spikes in performance. On the whole, the findings confirm that the suggested system enhances the speed of the response, as well as contributes to the reliability of web applications operations in the field.

VII. CONCLUSION

Web applications today produce huge volumes of logs which are essential in ensuring system reliability, performance and availability. Older log monitoring methods can be found to be reactive, rule based and cannot scale and cope with the size and complexity of the distributed systems. This paper has demonstrated an AI-based real-time log monitoring and incident response system on web applications to overcome

these weaknesses. The suggested system utilizes machine learning-based anomaly detection and smart event correlation to detect abnormal behavior of the system in a timely manner. The system generates timely alerts, insights and actionable incident data by analyzing live log streams across various components using an interactive dashboard to provide contextual information and timely alerts. This will cut down the overhead of operations, increase response time and general resilience of the system. Conclusively, the suggested solution contains a scalable and effective intelligent log monitoring and incident management framework applicable in current web settings. The current work is further developed by the implementation of the best deep learning models, the elaboration of cross-layer observability, and the incorporation of automated remediation strategies to make the system even more reliable and capable of self-healing.

VIII. REFERENCES

- [1] X. Zhao, Z. Jiang and J. Ma, "A Survey of Deep Anomaly Detection for System Logs," *2022 International Joint Conference on Neural Networks (IJCNN)*, Padua, Italy, 2022, pp. 1-8, doi: 10.1109/IJCNN55064.2022.9892726.
<https://ieeexplore.ieee.org/document/9892726>
- [2] M. Landauer, S. Onder, F. Skopik and M. Wurzenberger, "Deep learning for anomaly detection in log data: A survey," *Machine Learning with Applications*, vol. 12, p. 100470, 2023, doi:10.1016/j.mlwa.2023.100470.
<https://www.sciencedirect.com/science/article/pii/S2666827023000233>
- [3] S. Ramachandran, R. Agrahari, P. Mudgal, H. Bhilwaria, G. Long and A. Kumar, "Automated Log Classification Using Deep Learning," *Procedia Computer Science*, vol. 218, pp. 1722–1732, 2023, doi: 10.1016/j.procs.2023.01.150.
<https://www.sciencedirect.com/science/article/pii/S1877050923001503>
- [4] L. Wang, C. Zhang, R. Ding, Y. Xu, Q. Chen, W. Zou, Q. Chen, M. Zhang, X. Gao, H. Fan, S. Rajmohan, Q. Lin and D. Zhang, "Root Cause Analysis for Microservice Systems via Hierarchical Reinforcement Learning from Human Feedback," in *Proc. of the ACM International Conference*, 2023, doi: 10.1145/3580305.3599934.
<https://doi.org/10.1145/3580305.3599934>
- [5] T. Jia, Y. Wu, C. Hou and Y. Li, "LogFlash: Real-time Streaming Anomaly Detection and Diagnosis from System Logs for Large-scale Software Systems," *2021 IEEE 32nd International Symposium on Software Reliability Engineering (ISSRE)*, Wuhan, China, 2021, pp. 80-90, doi: 10.1109/ISSRE52982.2021.00021.
<https://ieeexplore.ieee.org/document/9700306>
- [6] A. Horváth *et al.*, "Anomaly Detection Algorithms for Real-Time Log Data Analysis at Scale," in *IEEE Access*, vol. 13, pp. 136288-136311, 2025, doi: 10.1109/ACCESS.2025.3594469.
<https://ieeexplore.ieee.org/document/11105402?denied=>
- [7] LogPAI, "LogHub: A comprehensive log dataset collection," GitHub repository, 2019. [Online]. Available: <https://github.com/logpai/loghub>. Accessed: Dec. 28, 2025.
- [8] Z. Jiang, J. Liu, J. Huang, Y. Li, Y. Huo, J. Gu, Z. Chen, J. Zhu and M. R. Lyu, "A Large-scale Evaluation for Log Parsing Techniques: How Far are We?," in *Proc. of the ACM SIGSOFT International Symposium on Software Testing and Analysis (ISSTA)*, 2024.
<https://arxiv.org/abs/2308.10828>
- [9] J. Zhu, S. He, P. He, J. Liu and M. R. Lyu, "LogHub: A Large Collection of System Log Datasets for AI-driven Log Analytics," in *Proc. of the IEEE International Symposium on Software Reliability Engineering (ISSRE)*, 2023.
<https://arxiv.org/abs/2008.06448>
- [10] L. Pham, H. Ha and H. Zhang, "Root Cause Analysis for Microservice System based on Causal Inference: How Far Are We?," in *Proc. of the Association for Computing Machinery (ACM) Conference*, 2024, doi: 10.1145/3691620.3695065.
<https://doi.org/10.1145/3691620.3695065>